

DynamoDBStreams Kinesis Adapter

V1->V2 Migration Notes

Andrew Muldowney - andrew.muldowney87@gmail.com

I tried my hand at converting the dynamodb-streams-kinesis-adapter from V1 to V2, AWS-SDK and Kinesis Client.

<https://github.com/amuldowney/dynamodb-streams-kinesis-adapter>

Java

```
<properties>
  <aws-java-sdkv2.version>2.17.242</aws-java-sdkv2.version>
  <software-kinesis.version>2.5.4</software-kinesis.version>
  <aws.dynamodblocal.version>[1.12,2.0)</aws.dynamodblocal.version>
</properties>
```

I was able to get all (existing) tests to pass which includes an integration test that writes some rows and reads them out of the stream. But during the process I had to either implement a `@KinesisClientInternalApi` or had to slow down the library in some way. I've noted the major ones here in this document. I hope this is helpful for others looking to do the same work and/or get some advice on ways to avoid implementing `@KinesisClientInternalApi` classes.

StreamsWorkerFactory.java

1. `ConfigsBuilder` doesn't take the `KinesisClientLibConfiguration` so we have to extract the specific configurations to override their settings.

Java

`StreamsWorkerFactory.java`

```
LeaseManagementConfig lmc = configsBuilder.leaseManagementConfig();
lmc =
lmc.initialLeaseTableReadCapacity(config.getInitialLeaseTableReadCapacity());
```

```
lmc =
lmc.initialLeaseTableWriteCapacity(config.getInitialLeaseTableWriteCapacity());
lmc = lmc.shardSyncIntervalMillis(config.getShardSyncIntervalMillis());
```

2. `DynamoDBLeaseManagementFactory` is an `@KinesisClientInternalApi` but we need to override it to set our own `DynamoDBStreamsShardSyncer` because there is nowhere else we can do it. Creating a `DynamoDBLeaseManagementFactory` is a metric ton of constructor args. There is a place to set a `hierarchicalShardSyncer` on the `LeaseManagementConfig` but it doesn't work.

Java

`StreamsWorkerFactory.java`

```
LeaseManagementConfig lmc = configsBuilder.leaseManagementConfig();
lmc = lmc.hierarchicalShardSyncer(shardSyncer); //DOESNT WORK

//Instead we must do the following
DynamoDBStreamsLeaseManagementFactory dynamoDBStreamsLeaseManagementFactory =
    new DynamoDBStreamsLeaseManagementFactory(
        streamsClient,
        dynamoDBClient,
        ...
        new DynamoDBStreamsShardSyncer(new StreamsLeaseCleanupValidator()),
        ...
        configsBuilder.retrievalConfig().streamTracker().isMultiStream() ?
            new DynamoDBMultiStreamLeaseSerializer() :
            new DynamoDBLeaseSerializer(),
        ...);
lmc.leaseManagementFactory(dynamoDBStreamsLeaseManagementFactory);
HierarchicalShardSyncer
```

3. `HierarchicalShardSyncer` is an `@KinesisClientInternalApi` but we need to override it as well to support some specific DDB needs.

Java

```
class DynamoDBStreamsShardSyncer extends HierarchicalShardSyncer...
```

* This `class` has been copied from `ShardSyncer` in `KinesisClientLibrary` and edited slightly to enable `DynamoDB Streams`

- * specific behavior. It is a helper `class` to sync leases with shards of the DynamoDB Stream.
- * It will create `new` leases/activities when it discovers `new` DynamoDB Streams shards (bootstrap/resharding).
- * It deletes leases `for` shards that have been trimmed from DynamoDB Stream.
- * It also ensures that leases `for` shards that have been completely processed are not deleted until their children
- * shards have also been completely processed.

`software.amazon.awssdk.services.kinesis.model.Record`

4. Thank goodness `SynchronousPrefetchingRetrievalFactory` exists, the default `RetrievalFactories` rely on stream connections to Kinesis which I didn't want to fake.

Java

```
RetrievalConfig rc = configsBuilder.retrievalConfig();
    rc = rc.retrievalFactory(new SynchronousPrefetchingRetrievalFactory(

rc.streamTracker().streamConfigList().get(0).streamIdentifier().streamName(),
    //always a single stream tracker for dynamodb streams
    streamsClient,
    new DynamoDBStreamsRecordsFetcherFactory(),
    1000,
    executorService,
    300,
    Duration.ofSeconds(30)
    ));
```

Model package

1. All the model work is now static instantiation / mapping since V2 data classes are final w/ builder. This is maybe a little slower than before. I've renamed all the classes `*Mapper` to indicate they aren't adapting but creating new objects.
2. `AmazonServiceExceptionTransformer` work is extensive, mapping the new types along with going from `editable` to `builders` all over

3. `RecordObjectMapper` & `RecordMapper` have changed dramatically. Since we can't sub-class `software.amazon.awssdk.services.kinesis.model.Record` (its final now) we have to actually store the DDB record in the Kinesis record bytes and pull it out on the other end.

Java

```
dynamoRecord = Record.builder()
    .awsRegion("us-east-1")
    .eventID(UUID.randomUUID().toString())
    .eventSource("aws:dynamodb")
    .eventVersion("1.1")
    .eventName(OperationType.MODIFY)
    .dynamodb(StreamRecord.builder()
        .approximateCreationDateTime(TEST_DATE)
        .keys(key)
        .oldImage(new HashMap<>(key))
        .newImage(new Image)
        .sizeBytes(Long.MAX_VALUE)
        .sequenceNumber(TEST_STRING)
        .streamViewType(StreamViewType.NEW_AND_OLD_IMAGES)
        .build()).build();

kinesisRecord = RecordMapper.convert(dynamoRecord);

@Test
public void testGetDataDeserialized() throws IOException {
    Whitebox.setInternalState(RecordMapper.class, ObjectMapper.class,
MAPPER);

    java.nio.ByteBuffer dynamoRecordAsData =
kinesisRecord.data().asByteBuffer();
    Record actual =
MAPPER.readValue(BinaryUtils.copyBytesFrom(dynamoRecordAsData),
Record.serializableBuilderClass()).build();
    assertEquals(dynamoRecord, actual);
}
```

ShardMapper

1. In `ShardMapper` we lie about the `HashKeyRange` to keep a different part of the code off our backs, parts of the KCL want to piece together hash keys which creates a loop that breaks the adapter.

Java

```
public class ShardMapper {

    static final BigInteger MIN_HASH_KEY = BigInteger.ZERO;
    static final BigInteger MAX_HASH_KEY = new
    BigInteger("2").pow(128).subtract(BigInteger.ONE);

    public static software.amazon.awssdk.services.kinesis.model.Shard
    convert(software.amazon.awssdk.services.dynamodb.model.Shard internalShard) {
        return software.amazon.awssdk.services.kinesis.model.Shard.builder()
            .shardId(internalShard.shardId())
            .parentShardId(internalShard.parentShardId())
            .adjacentParentShardId(null)

            .hashKeyRange(software.amazon.awssdk.services.kinesis.model.HashKeyRange.builder()
                .startingHashKey(MIN_HASH_KEY.toString())
                .endingHashKey(MAX_HASH_KEY.toString())
                .build())
            ...
    }
}
```